Claude

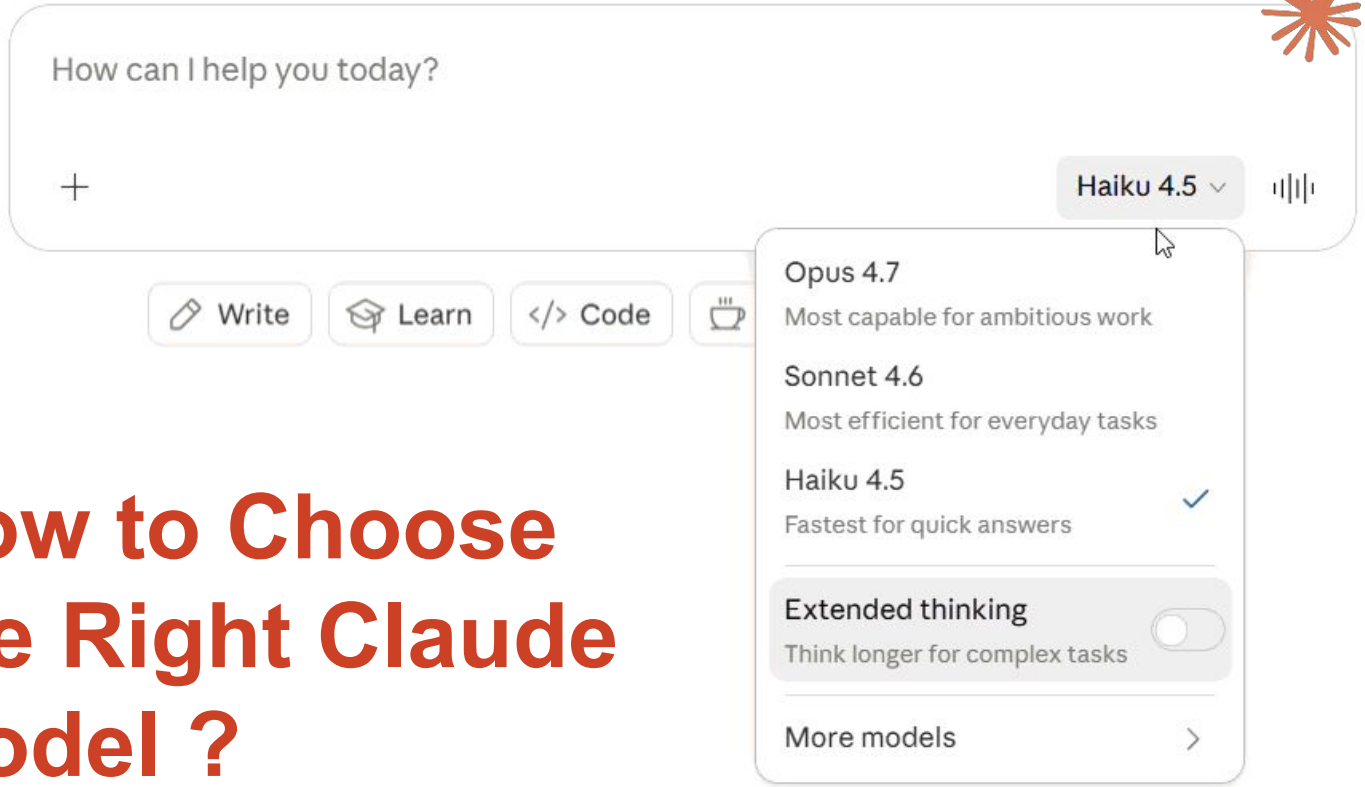

How to Choose the Right Claude Model ?

Claude

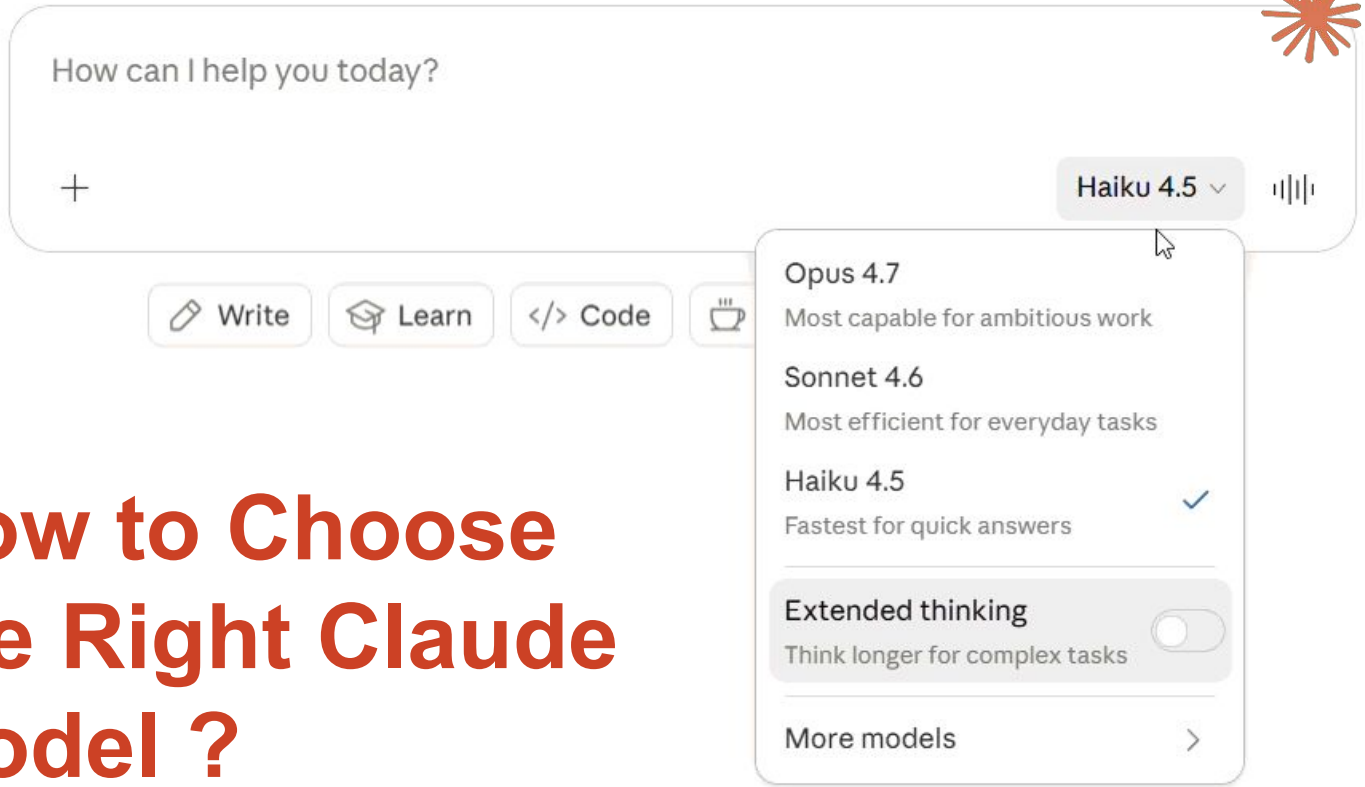

How to Choose the Right Claude Model ?

What will this video give you?

For Researchers

- Know each model's limits, backed by official system cards
- Balance cost, quality, and latency

For Developers

- Choose the right model for your use case before writing code
- Match inference techniques to model support — not every model supports every technique

For General Users

- Pick the right Claude model with confidence using just 4 questions
- See the clear trade-offs between latency, quality, and cost

Scope

- **Models covered:** the three latest Claude models available as of April 2026 — Opus 4.7, Sonnet 4.6, and Haiku 4.5
- **Inference techniques covered:** streaming, caching, batching, and thinking modes (extended, adaptive)

Ask yourself the following questions:

1. How complex is my task?
2. How long am I willing to wait for a response?
3. What is my budget?
4. What inference techniques do I want to use?

Ask yourself the following question …?

1.How complex is my task?

Complexity	Model	Examples of Tasks	Anthropic Description
Simple, repetitive, high-volume	**Haiku 4.5**	Grammar correction, Q&A chatbot, text classification, basic translation, data lookup, email auto-reply	"Fastest model with near-frontier intelligence" — designed for speed and high-volume tasks
Moderate — coding, analysis, reasoning	**Sonnet 4.6**	Code generation, document summarization, business reporting, math problem solving, financial analysis, web research	"Best combination of speed and intelligence" — sized for moderate-complexity professional tasks
Highest — advanced professional & engineering tasks	**Opus 4.7**	Complex code debugging, multi-step research, advanced software engineering, real-world professional consulting, scientific reasoning, long-context analysis	"Most capable model for complex reasoning and agentic coding" — designed for the highest-complexity tasks

Ask yourself the following question …?

2. How long am I willing to wait for a response?

Tolerance	Model	Anthropic's Description
Seconds — need it fast	**Haiku 4.5**	"The fastest model" — Comparative latency rated **Fastest**
Minutes — balanced	**Sonnet 4.6**	Comparative latency rated **Fast**; adaptive thinking lets the model regulate reasoning depth automatically
Longer — depth over speed	**Opus 4.7**	Comparative latency rated **Moderate**; uses adaptive thinking with effort levels (**low / medium / high / xhigh / max**) for fine-grained control over reasoning depth

Ask yourself the following question …?

3. What is my budget?

Budget	Model	Rationale	Official API Pricing (per 1M tokens)	Token Cost Concern
Low — high-volume, cost-sensitive	**Haiku 4.5**	Lowest per-token cost; well-suited for high-volume parallel tasks	Input: **$1.00** / Output: **$5.00** — Batch: $0.50/$2.50	
Mid — balance of quality and cost	**Sonnet 4.6**	Best combination of speed and intelligence" at mid-tier pricing	Input: **$3.00** / Output: **$15.00** — Batch: $1.50/$7.50	⚠️ **Thinking tokens** (extended or adaptive) are billed at the output rate ($15/M standard, $7.50/M batch). 💡 Use **adaptive thinking** so Claude only thinks when needed — skipping unnecessary reasoning on simple queries.
High — mission-critical output quality	**Opus 4.7**	Most capable general-access model; justified for complex professional tasks	Input: **$5.00** / Output: **$25.00** — Batch: $2.50/$12.50	🔴 **New tokenizer in Opus 4.7 may use up to 35% more tokens for the same text** compared to previous models

Ask yourself the following question …?

4. What inference techniques do I want to use? (API-level optimization)

Need	Inference Technique	Model Support	Notes
Real-time output for chatbots/agents	**Streaming**	All 3 models	Returns tokens as generated — better UX for live apps
Lower cost on repeated context	**Prompt Caching**	All 3 models	Cache hits cost 10% of input rate — huge savings on long prompts
50% discount, async OK	**Batch Processing**	All 3 models	Best for bulk jobs; not suitable for real-time apps
Smart auto-control of reasoning depth	**Adaptive Thinking**	Sonnet 4.6, Opus 4.7	Claude decides per-request; pairs with `effort` parameter
Predictable fixed thinking budget	**Extended Thinking**	Haiku 4.5, Sonnet 4.6	You set `budget_tokens`; **not supported on Opus 4.7**

Summary Decision Card

Decision Factor	Haiku 4.5	Sonnet 4.6	Opus 4.7
1. Task Complexity	Simple, repetitive, high-volume (grammar correction, Q&A, classification, translation)	Moderate (code generation, business reporting, document summarization, financial analysis)	Highest (complex debugging, multi-step research, advanced software engineering, professional consulting)
2. Response Wait Time	Seconds — fastest latency	Minutes — adaptive thinking self-regulates depth	Longer — designed for depth over speed
3. Budget (per 1M tokens)	Input: $1 / Output: $5 Batch: $0.50 / $2.50	Input: $3 / Output: $15 Batch: $1.50 / $7.50	Input: $5 / Output: $25 Batch: $2.50 / $12.50 ⚠️ New tokenizer may use up to 35% more tokens
4. Inference Techniques	Streaming, Prompt Caching, Batch Processing, **Extended Thinking**	Streaming, Prompt Caching, Batch Processing, **Extended Thinking, Adaptive Thinking**	Streaming, Prompt Caching, Batch Processing, **Adaptive Thinking** *(critical to offset tokenizer cost via caching)*

Conclusion

Ask yourself these 4 questions:

1. How complex is my task?
2. How long am I willing to wait for a response?
3. What is my budget?
4. What inference techniques do I want to use?

The answer depends on your situation. Pick the model that fits.

Haiku 4.5

Sonnet 4.6

Opus 4.7

Reference

Official Anthropic Documentation

- Anthropic document - https://platform.claude.com/docs/en/about-claude/models/overview
- Anthropic API Pricing — https://platform.claude.com/docs/en/about-claude/pricing
- Prompt Caching — https://platform.claude.com/docs/en/build-with-claude/prompt-caching
- Batch Processing — https://platform.claude.com/docs/en/build-with-claude/batch-processing
- Plans & Pricing — https://claude.com/pricing

*Content research and writing for this video assisted by Claude (Anthropic)

Happy Learning !